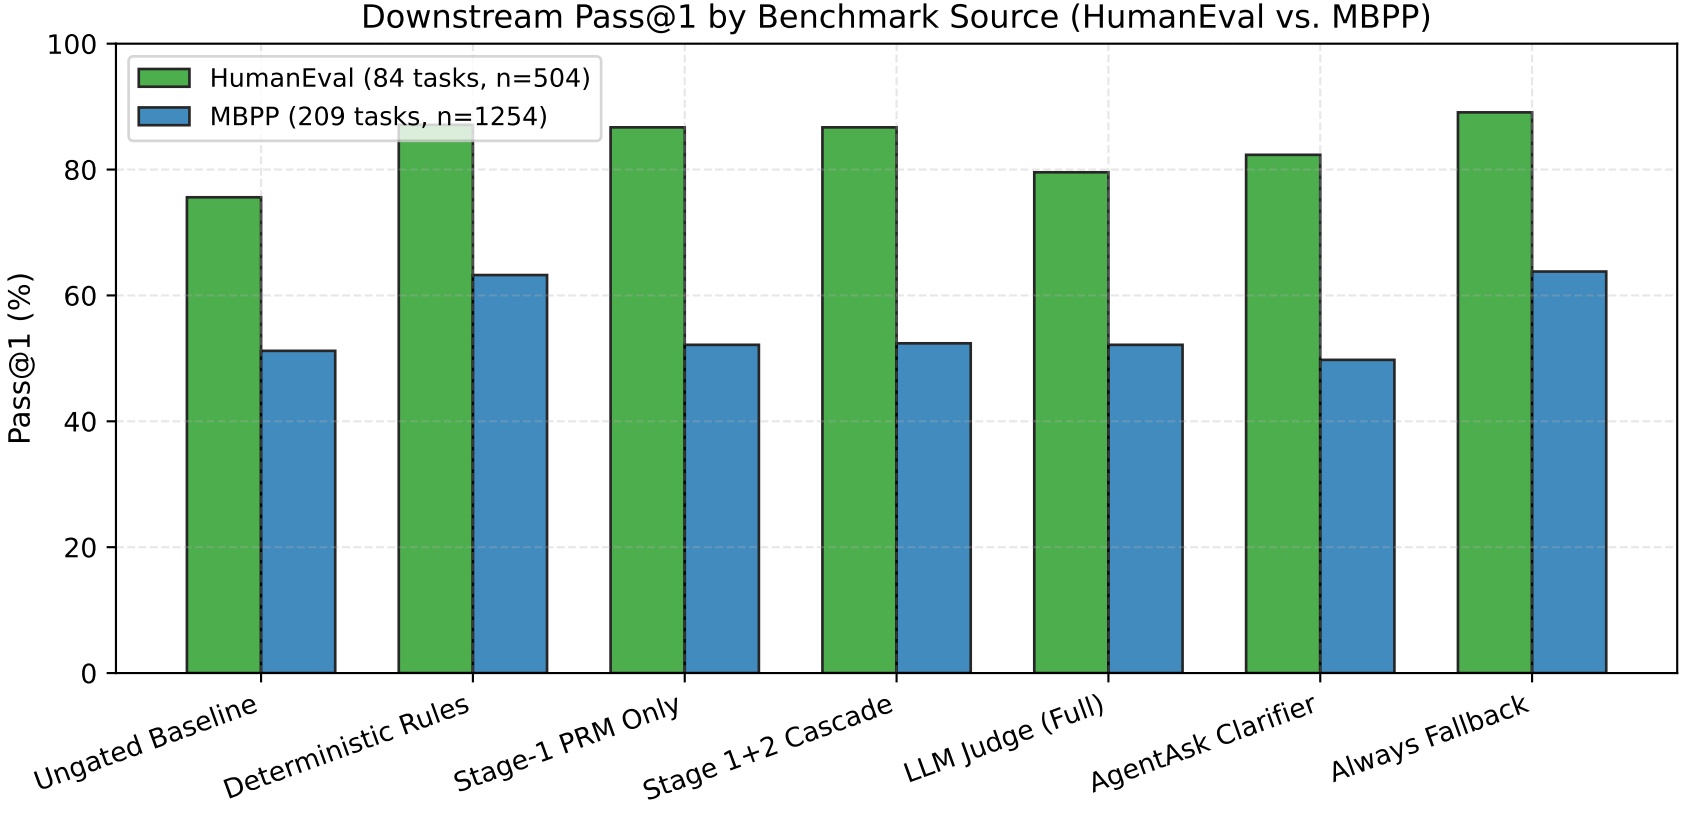
Downstream Pass@1 by Benchmark Source (HumanEval vs. MBPP)
HumanEval (84 tasks, n=504)
MBPP (209 tasks, n=1254)
Pass@1 (%)
100
80
60
40
20
0
Ungated Baseline
Deterministic Rules
Stage-1 PRM Only
Stage 1+2 Cascade
LLM Judge (Full)
AgentAsk Clarifier
Always Fallback